

Maximum Likelihood estimation

Jan Graffelman¹

¹Department of Statistics and Operations Research
Universitat Politècnica de Catalunya
Barcelona, Spain



UNIVERSITAT POLITÈCNICA
DE CATALUNYA
BARCELONATECH

jan.graffelman@upc.edu

September 19, 2022

Contents

- 1 Introduction
- 2 Model estimation
- 3 Maximum likelihood method
- 4 Other methods
- 5 Comparing estimators

Models

Scientists use **models** to describe and understand the phenomena they study.

We generally distinguish:

- Deterministic models.
- Stochastic models, also called **statistical models**.

Some examples:

- $V = I \times R$ (Ohm's law)
- $Y_i \sim N(\mu, \sigma^2)$ (Normal distribution)
- $Y_i = \beta_0 + \beta_1 X_i + \varepsilon_i$ (Linear regression model)
- $Y_t = \alpha + \beta Y_{t-1} + \varepsilon_t$ (Time series model)
- ...
- We will focus on statistical models, which have a probabilistic nature
- Probability theory is the foundation of statistics.

Population and sample: an opinion poll

We wish to know what percentage of all adult people in Spain favor legalization of Marijuana.

1000 Spanish adults are interviewed and their opinion is registered.

- Population: all adult people in Spain.
- Sample: 1000 Spanish adults.

Note: in statistics, populations are often assumed infinite.

6 / 49

- **Parameters** are fixed, unknown quantities that specify the population.
- Any number you compute using some sample of data is called a **statistic**.
- Statistics are random variables whereas parameters are not (unless you are a Bayesian).
- **Estimators** are statistics that are used to estimate the unknown parameters.

- $\hat{\mu} = \bar{x}$ with μ the true population mean.
- $\hat{\sigma} = s$ with σ the true population standard deviation.

- Let X_1, \dots, X_n be a random sample from a distribution $f(x|\theta_1, \dots, \theta_k)$.
- The **likelihood function** $L(\theta|\mathbf{x})$ is defined as

Example: Bernoulli distribution

Let X_1, \dots, X_n be a random sample with $X_i \sim \text{Bern}(p)$

$$P(X_1 = x_1 | p) = p^{x_1} (1 - p)^{1-x_1}$$

$$\begin{aligned} P(X_1 = x_1, \dots, X_n = x_n | p) &= \prod_{i=1}^n p^{x_i} (1-p)^{1-x_i} \\ &= p^{\sum_{i=1}^n x_i} (1-p)^{n-\sum_{i=1}^n x_i} \end{aligned}$$

$$L(p | x_1, \dots, x_n) = p^{\sum_{i=1}^n x_i} (1 - p)^{n - \sum_{i=1}^n x_i}$$

[illegible]

Let X_1, \dots, X_n be a random sample with $X \sim \exp(\mu)$.

$$f(x_i | \mu) = \frac{1}{\mu} \exp\left(\frac{-x_i}{\mu}\right)$$

$$f(x_1, \dots, x_n | \mu) = \frac{1}{\mu^n} \exp\left(-\frac{\sum_{i=1}^n x_i}{\mu}\right)$$

The likelihood function is:

$$L(\mu \mid x_1, \dots, x_n) = \frac{1}{\mu^n} \exp\left(\frac{-\sum_{i=1}^n x_i}{\mu}\right)$$

Example: the normal distribution

Let X_1, \dots, X_n be a random sample with $X \sim N(\mu, \sigma^2)$.

The joint density function is


$$f(x_1, \dots, x_n | \mu, \sigma^2) = \left(\frac{1}{\sqrt{2\pi\sigma^2}} \right)^n \exp \left\{ -\frac{1}{2\sigma^2} \sum_{i=1}^n (x_i - \mu)^2 \right\}$$

and the likelihood function is

$$L(\mu, \sigma^2 \mid x_1, \dots, x_n) = \left(\frac{1}{\sqrt{2\pi\sigma^2}} \right)^n \exp \left\{ -\frac{1}{2\sigma^2} \sum_{i=1}^n (x_i - \mu)^2 \right\}$$

Candidates for MLE

Some likelihood function



17 / 49

Example: exponential distribution

Let X be a random variable with $X \sim \exp(\lambda)$.

$$f(x | \lambda) = \lambda e^{-\lambda x}$$

We observe $x = 3$ (sample size $n = 1$)


$$L(\lambda | x = 3) = \lambda e^{-3\lambda}$$

$$L'(\lambda | x = 3) = e^{-3\lambda} (1 - 3\lambda)$$

$$\hat{\lambda} = \frac{1}{3} \quad L''(\lambda = 1/3 | x = 3) < 0$$

$$\lim_{\lambda \rightarrow 0} \lambda e^{-\lambda x} = 0$$

$$\lim_{\lambda \rightarrow \infty} \lambda e^{-\lambda x} = 0$$



Example: Bernoulli's distribution

Let X_1, \dots, X_n be a random sample with $X_i \sim \text{Bern}(p)$, and $\Theta = [0, 1]$.

$$L(p | x) = p^{\sum_{i=1}^n x_i} (1 - p)^{n - \sum_{i=1}^n x_i}$$

$$\log L(p | x) = \left(\sum_{i=1}^n x_i \right) \log p + \left(n - \sum_{i=1}^n x_i \right) \log(1 - p)$$

$$\frac{d}{dp} \log L(p | x) = \frac{\sum_{i=1}^n x_i}{p} - \frac{n - \sum_{i=1}^n x_i}{1 - p} = 0 \Leftrightarrow \hat{p} = \frac{\sum_{i=1}^n x_i}{n}$$

$\frac{\sum_{i=1}^n x_i}{n}$ is the only stationary point in $\Theta = [0, 1]$.


$$\begin{aligned} \frac{d^2}{dp^2} \log L(p | x) \Big|_{p=\hat{p}} &= -\frac{\sum_{i=1}^n x_i}{p^2} + \frac{\sum_{i=1}^n x_i - n}{(1 - p)^2} \Big|_{p=\hat{p}} = \\ &= -\frac{n\hat{p}}{\hat{p}^2} - \frac{n(1 - \hat{p})}{(1 - \hat{p})^2} = -\left(\frac{n}{\hat{p}} + \frac{n}{1 - \hat{p}} \right) < 0 \end{aligned}$$

Boundary points: $L(0 | x) = 0$ and $L(1 | x) = 0$

Example: Bernoulli distribution

Exercise (in R)

- Simulate 100 flips of a fair coin ($P(\text{"Heads"}) = P(\text{"Tail"}) = 0.50$)
- Calculate the value of the ML estimator, \hat{p}_{ML}
- Write a function that calculates the ML estimator as a function of p
- Make a plot of the likelihood function
- Verify graphically that \hat{p} maximizes the likelihood function



Example: normal distribution

Let X_1, \dots, X_n be a random sample with $X_i \sim N(\mu, 1)$.

$$L(\mu | \mathbf{x}) = \prod_{i=1}^n \frac{1}{\sqrt{2\pi}} e^{-\frac{1}{2}(x_i - \mu)^2} = \frac{1}{(2\pi)^{n/2}} e^{-\frac{1}{2} \sum_{i=1}^n (x_i - \mu)^2}$$

$$\ell(\mu, \mathbf{x}) = -\frac{n}{2} \ln(2\pi) - \frac{1}{2} \sum (x_i - \mu)^2$$

$$\frac{d}{d\mu} \ell(\mu | \mathbf{x}) = 0 \rightarrow \sum_{i=1}^n (x_i - \mu) = 0 \rightarrow \hat{\mu} = \bar{x}$$

$$\frac{d^2}{d\mu^2} \ell(\mu | \mathbf{x})|_{\mu=\bar{x}} < 0$$

$$\lim_{\mu \rightarrow +\infty} L(\mu | \mathbf{x}) = \lim_{\mu \rightarrow -\infty} L(\mu | \mathbf{x}) = 0$$

Additional examples

- Find the MLE for parameter λ of the Poisson distribution.
- Find the MLE for parameter p of the Geometric distribution.

Mutation rate in DNA

Sequence	Mutations
GACACGTATAAGGCATAACATACTGCGGTTTCGTTCCGATTATGAATCC...	—
GACACGTATAAGGCATAACATACTGCGGTTTCGTTCCGATTATGAATCC...	0
GACACGTATAAGGCATAACATACTGCGGTTTCGTTCCGATTATGAATCC...	0
GACACGTATAAGGCATAACATACTGCGGTTTCGTTCCGATTATGAATCC...	0
GACACGTATAAGGCATAACATACTGCGGTTTCGTTCCGATTATGAATCC...	1
GACACGTATAAGGCATAACATACTGCGGTTTCGTTCCGATTATGAATCC...	0
GACACGTATAAGGCATAACATACTGCGGTTTCGTTCCGATTATGAATCC...	1
GACACGTATAAGGCATAACATACTGCGGTTTCGTTCCGATTATGAATCC...	2
GACACGTATAAGGCATAACATACTGCGGTTTCGTTCCGATTATGAATCC...	0
⋮	

$$X = \text{Number of mutations} \sim \text{Pois}(\lambda)$$

Invariance property of the MLE

Let $\hat{\theta}$ be the MLE of θ . Then for any function $\tau(\theta)$ the MLE of $\tau(\theta)$ is $\tau(\hat{\theta})$.

Example:

- Let X_1, \dots, X_n be a random sample with $X_i \sim \text{Bern}(p)$.
- We wish to estimate $\ln\left(\frac{p}{1-p}\right)$ (the log odds).
- We know $\hat{p} = \frac{\sum_{i=1}^n x_i}{n}$
- The MLE of $\ln\left(\frac{p}{1-p}\right)$ is $\ln\left(\frac{\hat{p}}{1-\hat{p}}\right)$

Two parameters: the normal distribution

Let X_1, \dots, X_n be a random sample with $X_i \sim N(\theta, \sigma^2)$.

$$L(\theta, \sigma^2 | \mathbf{x}) = \frac{1}{(2\pi\sigma^2)^{n/2}} e^{-\frac{1}{2\sigma^2} \sum_{i=1}^n (x_i - \theta)^2}$$

$$\log L(\theta, \sigma^2 | \mathbf{x}) = -\frac{n}{2} \log 2\pi - \frac{n}{2} \log \sigma^2 - \frac{1}{2\sigma^2} \sum_{i=1}^n (x_i - \theta)^2$$

$$\frac{\partial}{\partial \theta} \log L(\theta, \sigma^2 | \mathbf{x}) = \frac{1}{\sigma^2} \sum_{i=1}^n (x_i - \theta) = 0$$

$$\frac{\partial}{\partial (\sigma^2)} \log L(\theta, \sigma^2 | \mathbf{x}) = -\frac{n}{2\sigma^2} + \frac{1}{2\sigma^4} \sum_{i=1}^n (x_i - \theta)^2 = 0$$

$$\hat{\theta} = \bar{x} \quad \text{and} \quad \hat{\sigma}^2 = \frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})^2$$

Two (and more) parameters: some general comments

- With two (or more) parameters, ML estimation amounts to the **maximization of a function that depends on two (or more) variables** (the parameters in this case).
- Such maximization in multiple variables is mathematically more involved, and explained in detail in the MDO course.
- In many advanced ML estimation problems, explicit (closed form) solutions do often not exist or are hard to find. In those cases, we **maximize the likelihood iteratively, with numerical methods**.
- The numerical methods typically require some **initial estimate or first guess** of the maximum.
- A sensible initial estimate can often be obtained by alternative estimation methods.

Two parameters: gamma distribution

Let X_1, \dots, X_n be a random sample with $X_i \sim \Gamma(\alpha, \lambda)$.

$$f(x | \alpha, \lambda) = \frac{1}{\Gamma(\alpha)} \lambda^\alpha x^{\alpha-1} e^{-\lambda x} \quad \text{for } 0 \leq x < \infty$$

$$L(\theta) = L(\alpha, \lambda) = \left(\frac{1}{\Gamma^n(\alpha)} \right) \lambda^{n\alpha} \left(\prod_{i=1}^n x_i \right)^{\alpha-1} e^{-\lambda \sum_{i=1}^n x_i}$$

$$\ell(\theta) = \log L(\theta) = -n \log \Gamma(\alpha) + n \alpha \log \lambda + (\alpha - 1) \sum_{i=1}^n \log x_i - \lambda \sum_{i=1}^n x_i$$

$$\frac{\partial \ell}{\partial \alpha} = -n \frac{\Gamma'(\alpha)}{\Gamma(\alpha)} + n \log \lambda + \sum_{i=1}^n \log x_i$$

$$\frac{\partial \ell}{\partial \lambda} = n \alpha \frac{1}{\lambda} - \sum_{i=1}^n x_i$$

$$n \alpha \frac{1}{\lambda} - \sum_{i=1}^n x_i = 0 \Leftrightarrow \hat{\lambda} = \frac{n \hat{\alpha}}{\sum_{i=1}^n x_i} = \frac{\hat{\alpha}}{\bar{x}_n}$$

$$-n \frac{\Gamma'(\hat{\alpha})}{\Gamma(\hat{\alpha})} + n \log \frac{\hat{\alpha}}{\bar{x}_n} + \sum_{i=1}^n \log x_i = 0$$

There is no explicit solution.

The Newton-Raphson method

- We compute $\hat{\alpha}$ iteratively (Newton-Raphson method)
- Roots of the function $f(\alpha) = 0$ can be found by:

$$\hat{\alpha}_{n+1} = \hat{\alpha}_n + h_n \quad h_n = -\frac{f(\hat{\alpha}_n)}{f'(\hat{\alpha}_n)}$$

- For our problem:

$$f(\hat{\alpha}) = -n \frac{\Gamma'(\hat{\alpha})}{\Gamma(\hat{\alpha})} + n \log \frac{\hat{\alpha}}{\bar{x}} + \sum_{i=1}^n \log x_i = -n \frac{\Gamma'(\hat{\alpha})}{\Gamma(\hat{\alpha})} + n \log \hat{\alpha} - n \log \bar{x} + \sum_{i=1}^n \log x_i$$

and

$$f'(\hat{\alpha}) = -n \frac{d}{d\hat{\alpha}} \left(\frac{\Gamma'(\hat{\alpha})}{\Gamma(\hat{\alpha})} \right) + \frac{n}{\hat{\alpha}},$$

- The fraction $\frac{\Gamma'(\hat{\alpha})}{\Gamma(\hat{\alpha})}$ is known as the **digamma** function.
- Its derivative $\frac{d}{d\hat{\alpha}} \left(\frac{\Gamma'(\hat{\alpha})}{\Gamma(\hat{\alpha})} \right)$ is known as the **trigamma** function.
- An initial value α_0 is needed. We could use $\alpha_0 = 1$ or take the value of the estimator obtained by the method of moments.

Two parameters: gamma distribution

Consider a sample of 10.000 observations of a $\Gamma(\alpha = 2, \lambda = 3)$ distribution. The mean of the sample is 0.659269. We find the value of the MLE iteratively

i	α	$f(\alpha)$	$f'(\alpha)$	h
0	1.000000	3.102326e+03	-6449.341	4.810300e-01
1	1.481030	1.071487e+03	-2755.763	3.888170e-01
2	1.869847	2.364494e+02	-1672.638	1.413632e-01
3	2.011210	1.764700e+01	-1432.216	1.232146e-02
4	2.023532	1.141968e-01	-1413.740	8.077639e-05
5	2.023612	4.844877e-06	-1413.620	3.427285e-09
6	2.023612	-1.818989e-12	-1413.620	-1.286760e-15

$\hat{\alpha} = 2.023612$. Using $\hat{\lambda} = \frac{\hat{\alpha}}{\bar{x}}$ we find $\hat{\lambda} = \frac{2.023612}{0.659269} = 3.069479$.
By the method of moments, we find:

$$\hat{\alpha}_{MM} = \frac{\bar{x}^2}{\frac{1}{n}(\sum_{i=1}^n x_i^2) - \bar{x}^2} = 2.045233$$

this is a better initial point, from which we converge faster to the maximum.

i	α	$f(\alpha)$	$f'(\alpha)$	h
0	2.045233	-3.022058e+01	-1382.049	-2.186650e-02
1	2.023367	3.471806e-01	-1413.984	2.455335e-04
2	2.023612	4.477239e-05	-1413.620	3.167216e-08
3	2.023612	-1.818989e-12	-1413.620	-1.286760e-15

Some special cases

- Sometimes the support of the density depends on the parameter of interest.
- It then makes sense to use indicator variables that account for this.
- Examples: uniform distribution, distributions with a translation parameter, ...

Example: uniform distribution


$$X \sim U[0, \theta]$$

$$f(x, \theta) = \frac{1}{\theta} \cdot \mathbb{I}_{0 \leq x \leq \theta}$$

$$\begin{aligned} L(\theta | x_1, x_2, \dots, x_n) &= \prod_{i=1}^n \frac{1}{\theta} \prod_{i=1}^n \mathbb{I}_{x_i \leq \theta} \\ &= \frac{1}{\theta^n} \mathbb{I}_{x_{(n)} \leq \theta} \\ &= \frac{1}{\theta^n} \mathbb{I}_{\theta \geq x_{(n)}} \end{aligned}$$

$$\hat{\theta}_{ML} = X_{(n)}$$

Likelihood function



Exercise:

$$X \sim U(\alpha, 1) \quad f(x) = \frac{1}{1 - \alpha} \quad 0 < \alpha < x < 1$$

Find the ML estimator for α

Precision of the ML estimator

- A point estimate obtained by ML is, by itself, not very informative.
- We need to specify its **precision**.
- The precision depends on the **variance** or the **Fisher information** of the ML estimator.

Fisher information of a sample


Let X_1, \dots, X_n be a random sample with

$$f(\mathbf{x} \mid \theta) = \prod_{i=1}^n f(x_i \mid \theta)$$

The **Fisher information** about θ contained in \mathbf{x} is defined by

$$I_{\mathbf{x}}(\theta) = E_{\theta} \left[\left(\frac{\partial}{\partial \theta} \ln(f(\mathbf{x} \mid \theta)) \right)^2 \right]$$

Interpretation of Fisher Information



Fisher information relates to the **curvature** of the likelihood function

Cramér-Rao lower bound

- For any unbiased estimator ($E(\hat{\theta}) = \theta$), there exists a lower bound on its variance.
- This bound equals the reciprocal of the Fisher information.

$$V(\hat{\theta}) \geq \frac{1}{I_{\mathbf{x}}(\theta)}$$

- An unbiased estimator that attains the Cramér-Rao lower bound is called **efficient**.

Asymptotic distribution of the ML estimator

Let X_1, \dots, X_n be i.i.d. with density $f(x|\theta)$, and let $\hat{\theta}$ be the MLE of θ . Under regularity conditions we have

$$\hat{\theta}_n \rightarrow N\left(\theta, \frac{1}{I_x(\theta)}\right)$$

where $1/I_x(\theta)$ is the Cramér-Rao lower bound.

Thus, MLE are asymptotically (for large samples)

- unbiased,
- efficient,
- and normally distributed.

Interval estimation with maximum likelihood estimators

- Having the variance and the distribution of the ML estimator, we can now say something about uncertainty.
- A **confidence interval** is an expression of the uncertainty of the estimate.
- A classical result, with $X_i \sim N(\mu, \sigma^2)$, is

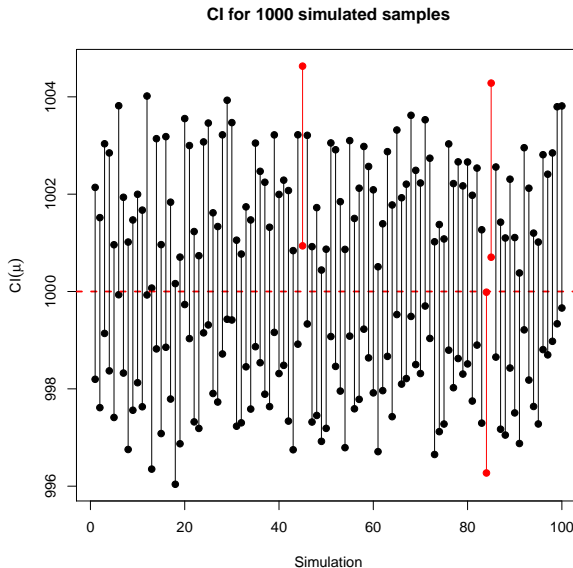
$$CI(\mu)_{1-\alpha} = \bar{X} \pm z_{\alpha/2} \frac{\sigma}{\sqrt{n}} \quad (1)$$

where $\bar{X} = \hat{\mu}_{ML}$, and $\frac{\sigma}{\sqrt{n}} = \sqrt{\frac{\sigma^2}{n}} = \sqrt{V(\hat{\mu})}$.

- Term $\frac{\sigma}{\sqrt{n}}$ (σ estimated by s) is called the **standard error of the mean**.
- Term $z_{\alpha/2} \frac{\sigma}{\sqrt{n}} \approx 2 \frac{\sigma}{\sqrt{n}}$ when $\alpha = 0.05$ is the **error margin**.
- Equation (1) holds in general for ML estimators:



$$CI(\theta)_{1-\alpha} = \hat{\theta} \pm z_{\alpha/2} \sqrt{V(\hat{\theta})} \quad (2)$$

Frequentist interpretation of a confidence interval



A practical example of ML estimation

In a study on physical characteristics of students, data on the weight (in kg) of $n = 91$ students is collected.




$$\hat{\mu}_{ML} = \bar{X} = 70.58$$

$$CI(\mu)_{0.95} = (68.65, 72.52)$$

Any problem?

Stratifying by gender



R instructions

```
X <- read.table(
  "http://www-eio.upc.es/~jan/data/StudentWeight.txt",
  header=TRUE)
```

```
weight <- X[,2]
sex <- X[,4]
```

```
hist(weight,breaks=12)
boxplot(weight)
```

```
boxplot(weight~sex,names=c("Males","Females"))
```


```
mean(weight[sex==0])
mean(weight[sex==1])
```

```
t.test(weight[sex==0]) % for obtaining the confidence
t.test(weight[sex==1]) % intervals
```

$$\hat{\mu}_{\text{males}} = \bar{X} = 72.62 \quad CI(\mu)_{0.95} = (70.73, 74.52)$$

$$\hat{\mu}_{\text{females}} = \bar{X} = 59.36 \quad CI(\mu)_{0.95} = (56.23, 62.48)$$

Example: ML estimation of the rate of an exponential distribution



- What is the rate of decay?
- What is the precision of a rate estimate?

```
> fitdistr(x,"exponential")
      rate
0.498116487
(0.004981165)
```

Density and likelihood:

$$f(x|\lambda) = \lambda e^{-\lambda x} \quad L(\lambda|\mathbf{x}) = \lambda^n e^{-\lambda \sum x_i}$$

With some algebra, it follows that

$$\hat{\lambda} = 1/\bar{x}, \quad I_n(\lambda) = n/\lambda^2 \quad V(\hat{\lambda}) = \lambda^2/n$$

$$CI_{1-\alpha}(\lambda) = \hat{\lambda} \pm z_{\alpha/2} \sqrt{V(\hat{\lambda})} = \hat{\lambda} \pm z_{\alpha/2} \frac{\hat{\lambda}}{\sqrt{n}}$$

Descriptive statistics of a sample of $n = 10.000$ waiting times

	N	N*	Mean	Stdev	Med	Q1	Q3	Min	Max
X	10000	0	2.0075	2	1.397	0.579	2.768	0.001	18.163

$$\hat{\lambda} = 1/2.0075 = 0.49812$$

$$CI_{0.95}(\lambda) = 0.49812 \pm 1.96 \frac{0.49812}{\sqrt{10000}} = (0.4884; 0.5079)$$

- The method of moment
- Bayesian methods

Moments

- $$V(X) = E(X - E(X))^2$$

Method of moments

Sample	Population
$m_1 = \frac{1}{n} \sum_{i=1}^n X_i$	$\mu_1 = E(X)$
$m_2 = \frac{1}{n} \sum_{i=1}^n X_i^2$	$\mu_2 = E(X^2)$
$m_3 = \frac{1}{n} \sum_{i=1}^n X_i^3$	$\mu_3 = E(X^3)$
\vdots	\vdots
$m_k = \frac{1}{n} \sum_{i=1}^n X_i^k$	$\mu_k = E(X^k)$

- Equate sample moments to population moments.
- Use as many moments as the number of parameters you need to estimate.
- Write the parameters as a function of the sample moments.

Method of moments

Let $X \sim U(0, \theta)$. We take a simple random sample of size n .

$$m_1 = \bar{X} \quad E(X) = \frac{\theta}{2}$$

$$\frac{\hat{\theta}_{MM}}{2} = \bar{X} \rightarrow \hat{\theta}_{MM} = 2\bar{X}$$

Exercise:

Let $X \sim \text{Exp}(\lambda)$. We take a simple random sample of size n .

$$f(x, \lambda) = \lambda e^{-\lambda x}$$

Find an estimator $\hat{\lambda}_{MM}$ for λ by using the method of moments.

Method of moments (Normal distribution)

Let X_1, X_2, \dots, X_n be random sample of size n from a $N(\mu, \sigma^2)$ distribution.

$$\mu_1 = E(X) = \mu \quad \mu_2 = E(X^2) = V(X) + E(X)^2 = \sigma^2 + \mu^2$$

$$m_1 = \bar{X} \quad m_2 = \frac{1}{n} \sum_{i=1}^n X_i^2$$

$$\hat{\mu}_{MM} = \bar{X}$$

$$\hat{\sigma}^2 + \hat{\mu}^2 = \frac{1}{n} \sum_{i=1}^n X_i^2$$

$$\hat{\sigma}^2 = \frac{1}{n} \sum_{i=1}^n X_i^2 - \bar{X}^2 = \frac{1}{n} \sum_{i=1}^n (X_i - \bar{X})^2$$

The Bayesian approach

- In classical, frequentist statistics, θ is assumed to be an unknown, fixed quantity.
- In the Bayesian approach, θ is a random variable, and its variation is described by a distribution, the **prior distribution**, $\pi(\theta)$.
- The **prior distribution**, $\pi(\theta)$, is subjective, and chosen by the investigator.
- A sample X_1, X_2, \dots, X_n is observed, and in the light of this data the distribution of θ is updated.
- The newly obtained distribution is called the **posterior distribution**, $\pi(\theta|\mathbf{x})$.
- The posterior distribution is

$$\pi(\theta|\mathbf{x}) = \frac{f(\theta, \mathbf{x})}{m(\mathbf{x})} = \frac{f(\mathbf{x}|\theta)\pi(\theta)}{m(\mathbf{x})}$$


with

$$m(\mathbf{x}) = \int f(\mathbf{x}|\theta)\pi(\theta)d\theta$$

- A **point estimate for θ** is obtained by calculating the **expectation (or the median) of the posterior distribution**.
- The **posterior distribution is proportional to the likelihood function and the prior**

$$\pi(\theta|\mathbf{x}) \propto f(\mathbf{x}|\theta)\pi(\theta)$$

Some criteria for comparing estimators



- Bias = $E(\hat{\theta}) - \theta$
- Variance $V(\hat{\theta})$ (or Precision = $\frac{1}{V(\hat{\theta})}$)
- Mean squared error $MSE(\hat{\theta}) \equiv E((\hat{\theta} - \theta)^2) = V(\hat{\theta}) + (Bias(\hat{\theta}))^2$

